Homework 4

Academic Honesty

Aside from the narrow exception for collaboration on homework, all work submitted in this course must be your own. Cheating and plagiarism will not be tolerated. If you have any questions about a specific case, please ask me. We will be checking for this!

NYU Poly's Policy on Academic Misconduct:
http://engineering.nyu.edu/academics/code-of-conduct/academic-misconduct

General Notes:

- Read the assignment carefully, including what files to include.
- Don't assume limitations unless they are explicitly stated.
- Treat provided examples as just that, not exhaustive list of cases that should work.
- When in doubt regarding what needs to be done, ask. Another option is test it in the real UNIX operating system. Does it behave the same way?
- **TEST** your solutions, make sure they work. It's obvious when you didn't test the code.

Parallel Hashtable with pthreads

In this assignment, you will take a non thread-safe version of a hash table and modify it so that it correctly supports running with multiple threads. This ****does not involve xv6****; xv6 doesn't currently support multiple threads of execution, and while it is possible to do parallel programming with processes, it's tricky to arrange access to some shared resource. Instead you will do this assignment on a multicore machine. Essentially any desktop machine made in the past 7 years or so should have multiple cores.

Start by downloading the attached file, **parallel_hashtable.c** to your local machine and you will compile it with the following command:

```
$ gcc -pthread parallel_hashtable.c -o parallel_hashtable
```

Now run it with one thread:

```
$ ./parallel_hashtable 1
[main] Inserted 100000 keys in 0.006545 seconds
[thread 0] 0 keys lost!
[main] Retrieved 100000/100000 keys in 4.028568 seconds
```

So with one thread the program is correct. But now try it with more than one thread:

```
$ ./parallel_hashtable 8
[main] Inserted 100000 keys in 0.002476 seconds
[thread 7] 4304 keys lost!
[thread 6] 4464 keys lost!
[thread 2] 4273 keys lost!
[thread 1] 3864 keys lost!
[thread 4] 4085 keys lost!
[thread 5] 4391 keys lost!
[thread 3] 4554 keys lost!
[thread 0] 4431 keys lost!
[main] Retrieved 65634/100000 keys in 0.792488 seconds
```

Play around with the number of threads. You should see that, in general, the program gets faster as you add more threads up until a certain point. However, sometimes items that get added to the hash table get lost.

Part 1

Find out under what circumstances entries can get lost. Update `parallel_hashtable.c` so that `insert` and `retrieve` do not lose items when run from multiple threads. Verify that you can now run multiple threads without losing any keys. Compare the speedup of multiple threads to the version that uses no mutex -- you should see that there is some overhead to adding a mutex.

You will probably need:

```
pthread_mutex_t lock;              // declare a lock
pthread_mutex_init(&lock, NULL);   // initialize the lock
pthread_mutex_lock(&lock);         // acquire lock
pthread_mutex_unlock(&lock);       // release lock
```

We will see this API in class this week. You can also use `man` to get more documentation on any of these.

Once you have a solution to this problem save it to a file called **'parallel_mutex.c'**

2

Part 2

For part 2 continue working with **'parallel_mutex.c' **
Does **retrieving** an item from the hash table require a lock? Update the code so that **multiple** `retrieve` operations can run in parallel.

Part 3

For part 4 continue working with **'parallel_mutex.c'**
Update the code so that some `insert` operations can run in parallel.

Hint: if two inserts are being done on *different* buckets in parallel, is a lock needed? What are the shared resources that need to be protected by mutexes?

Submitting

Upload **'parallel_mutex.**c to NYU Classes.